

Introduction:

In this project, we used IBM Attrition and Occupancy detection datasets to implement ANN and KNN algorithm. We used caret package to implement both the algorithms.

Dataset Information:

IBM Attrition:

This is a fictional dataset created by IBM data scientists. The dataset contains variables such as monthly income, overtime hours, tenure etc. We must develop a model to understand whether the employee will leave the organization or not.

This problem is interesting as it helps us understand what are the important variables that make an employee leave an organization.

Occupancy Detection:

Problem statement: Experimental data is used to predict room occupancy (binary classification) from various parameters such as Temperature, Light, and CO2

The dataset helps us in predicting if a room is occupied or not, depending on the room conditions

Data Preparation:

IBM Attrition:

Performed the following operations on the dataset before designing the model with ANN and KNN algorithms.

- The dataset is cleaned so omitting missing values didn't change the row count
- Dropped the categorical variables with only 1 just level
- Converted categorical variables into dummy variables
- Scale the entire dataset using normalize function
- Removed class imbalance using smote methods.
- Partitioned data into 50:50 ratio of Train and Test Dataset

Occupancy Detection:

Performed the following operations on the dataset before designing the model with ANN and KNN algorithms.

- Converted the continuous variable Occupancy into a factor variable with 2 levels
- Normalized all the predictor columns (Humidity, Temperature, CO2, Light, Humidity Ratio)
- Partitioned data into 50:50 ratio of Train and Test Dataset

Algorithm Implementation:

Artificial Neural Network (ANN) on IBM Attrition Dataset:

Trained a model with one hidden layer and 10 nodes, and weight 0.1 using caret package and then tested against Test Dataset. Confusion Matrix for both models are as follows:

For Train Dataset
Confusion Matrix and Statistics

	Reference	
Prediction	No	Yes
No	747	1
Yes	0	663

Accuracy : 0.9993
 95% CI : (0.9961, 1)
 No Information Rate : 0.5294
 P-Value [Acc > NIR] : <2e-16

 Kappa : 0.9986
 McNemar's Test P-Value : 1

 Sensitivity : 1.0000
 Specificity : 0.9985
 Pos Pred Value : 0.9987
 Neg Pred Value : 1.0000
 Prevalence : 0.5294
 Detection Rate : 0.5294
 Detection Prevalence : 0.5301
 Balanced Accuracy : 0.9992

 'Positive' Class : No

For Test Dataset
Confusion Matrix and Statistics

	Reference	
Prediction	No	Yes
No	292	19
Yes	27	265

Accuracy : 0.9237
 95% CI : (0.8996, 0.9436)
 No Information Rate : 0.529
 P-Value [Acc > NIR] : <2e-16

 Kappa : 0.8471
 McNemar's Test P-Value : 0.302

 Sensitivity : 0.9154
 Specificity : 0.9331
 Pos Pred Value : 0.9389
 Neg Pred Value : 0.9075
 Prevalence : 0.5290
 Detection Rate : 0.4842
 Detection Prevalence : 0.5158
 Balanced Accuracy : 0.9242

 'Positive' Class : No

For Train Dataset: The model accuracy is pretty good (99%) for training dataset. From the confusion Matrix, No Type 1 error and one record were classified incorrectly. Model sensitivity is 1. AUC is 1


For Test Dataset: The model accuracy for Test Dataset is 92% with sensitivity and specificity 91% and 93% respectively. The Kappa is 0.84. AUC is 0.89

To understand the impact of nodes, hidden layers, and an activation function, we did some experimentation with these parameters and regenerated the models.

Model accuracy with Sigmoid activation function was 1 whereas for tanh activation function was 0.99.

The Model was regenerated with 2 hidden layers with nodes 15 and 10 each and initial weight decay reduced to 0.1 from 0.5. The Model AUC is 1. The newly generated model was tested against Test dataset. Model accuracy for Test dataset increased from 92% to 93%.

The ROC curves for models are as shown below which indicate the area under the curve for train model, test model before experimentation and after experimentation. The ROC curve is as shown below.



Artificial Neural Network (ANN) on Occupancy Dataset:

Trained a model with one hidden layer and 10 nodes, and weight 0.1 using caret package and then tested against Test Dataset. Confusion Matrix for both models are as follows:

For Train Dataset: The model accuracy is pretty good (98%) for training dataset. From the confusion Matrix, we have 2 observations of Type 1 error and 2 observations of type 2 error. Model sensitivity is 0.9744.

For Test Dataset: The model accuracy for Test Dataset is 94% with sensitivity and specificity 93% and 100% respectively. The Kappa is 0.7953.

To understand the impact of nodes, hidden layers, and an activation function, we did some experimentation with these parameters and regenerated the models.

For Train Dataset
> print(conf1)
Confusion Matrix and Statistics

	Reference	
Prediction	yes	no
yes	2855	2
no	75	1139

Accuracy : 0.9811
95% CI : (0.9764, 0.985)
No Information Rate : 0.7197
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.954
McNemar's Test P-Value : 2.303e-16

Sensitivity : 0.9744
Specificity : 0.9982
Pos Pred Value : 0.9993
Neg Pred Value : 0.9382
Prevalence : 0.7197
Detection Rate : 0.7013
Detection Prevalence : 0.7018
Balanced Accuracy : 0.9863

'Positive' Class : yes

For Test Dataset
> print(conf2)
Confusion Matrix and Statistics

	Reference	
Prediction	yes	no
yes	3243	0
no	241	588


Accuracy : 0.9408
95% CI : (0.9331, 0.9479)
No Information Rate : 0.8556
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.7953
McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.9308
Specificity : 1.0000
Pos Pred Value : 1.0000
Neg Pred Value : 0.7093
Prevalence : 0.8556
Detection Rate : 0.7964
Detection Prevalence : 0.7964
Balanced Accuracy : 0.9654

'Positive' Class : yes

The ROC curves for models are as shown below which indicate the area under the curve for train model, test model before experimentation. The ROC curve is as shown below.

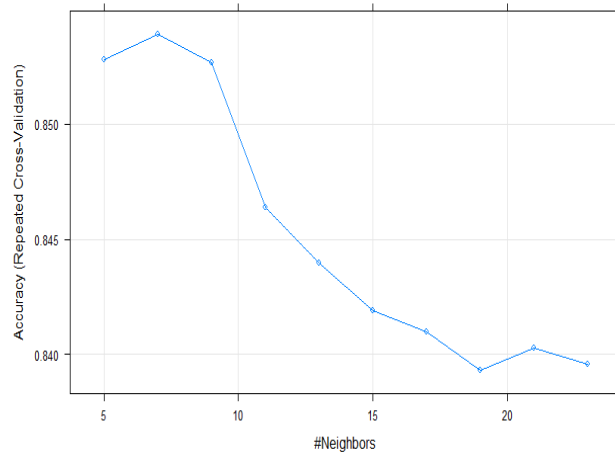


KNN Algorithm

K Nearest Neighbors (KNN) on IBM Attrition Dataset:

Trained a K nearest model on 50% data and tested the model on the rest 50% of the data.

Used repeated cross-validation to avoid overfitting on training data and reduce variance.



Confusion Matrix and Statistics

Prediction	Reference	
	No	Yes
No	609	106
Yes	8	12

Accuracy : 0.8449
95% CI : (0.8167, 0.8703)
No Information Rate : 0.8395
P-value [Acc > NIR] : 0.3663


Kappa : 0.1336
McNemar's Test P-value : <2e-16

Sensitivity : 0.9870
Specificity : 0.1017
Pos Pred Value : 0.8517
Neg Pred Value : 0.6000
Prevalence : 0.8395
Detection Rate : 0.8286
Detection Prevalence : 0.9728
Balanced Accuracy : 0.5444

'Positive' Class : No

Using repeated cross-validation, the best accuracy was achieved with K = 7.

The confusion matrix shows an accuracy of 84.49% but the Specificity of the model is low at .1017

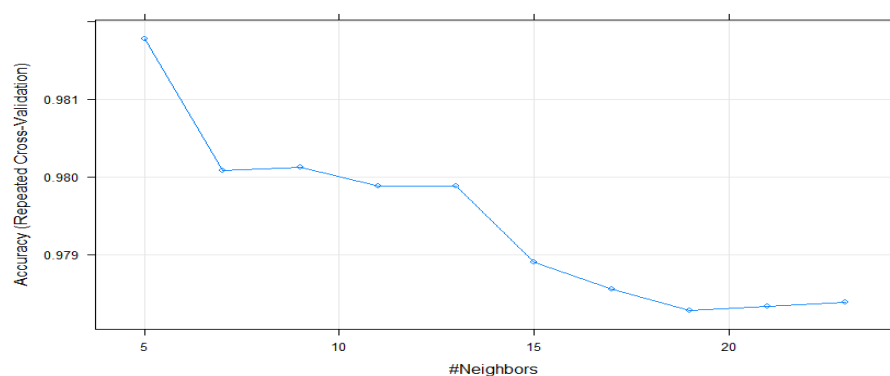


The AUC score of the model is also very low at 0.5443645

K Nearest Neighbors (KNN) on Occupancy Detection Dataset:

Trained a K nearest model on 50% data and tested the model on the rest 50% of the data.

Used repeated cross-validation to avoid overfitting on training data and reduce variance.



k-Nearest Neighbors

4071 samples
5 predictor
2 classes: '0', '1'

No pre-processing

Resampling: Cross-validated (10 fold, repeated 10 times)

Summary of sample sizes: 3664, 3664, 3664, 3664, 3664, ...

Resampling results across tuning parameters:

k	Accuracy	Kappa
5	0.9817732	0.9555820
7	0.9800781	0.9514888
9	0.9801272	0.9516735
11	0.9798817	0.9511274
13	0.9798816	0.9511916
15	0.9788990	0.9488368
17	0.9785550	0.9480397
19	0.9782848	0.9474026
21	0.9783339	0.9475247
23	0.9783830	0.9476461

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was k = 5.

k-Nearest Neighbors

735 samples
85 predictor
2 classes: 'No', 'Yes'

No pre-processing

Resampling: Cross-validated (10 fold, repeated 10 times)

Summary of sample sizes: 661, 662, 661, 661, 661, 662, ...

Resampling results across tuning parameters:

k	Accuracy	Kappa
5	0.8527914	0.19087514
7	0.8538742	0.16243935
9	0.8526580	0.14367416
11	0.8463973	0.08359400
13	0.8439556	0.06027367
15	0.8419194	0.04343927
17	0.8409622	0.03619915
19	0.8393314	0.01746633
21	0.8402755	0.02126545
23	0.8395998	0.01465638

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was k = 7.

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	3464	4
1	20	584

Accuracy : 0.9941
95% CI : (0.9912, 0.9962)
No Information Rate : 0.8556
P-value [Acc > NIR] : <2e-16


Kappa : 0.9764
McNemar's Test P-value : 0.0022

Sensitivity : 0.9943
Specificity : 0.9932
Pos Pred Value : 0.9988
Neg Pred Value : 0.9669
Prevalence : 0.8556
Detection Rate : 0.8507
Detection Prevalence : 0.8517
Balanced Accuracy : 0.9937

'Positive' Class : 0

Using repeated cross-validation, the best accuracy was achieved with K = 5.

The confusion matrix shows an accuracy of 99.41%, both the Sensitivity and Specificity are high at 0.9943 and 0.9932 respectively.



The AUC score of the model is very high at 0.9937284.

Comparison: Based on ROC and AUC for both the algorithms, we can clearly observe ANN is doing well as compare to KNN. KNN performance is better for occupancy dataset but not for IBM attrition.

Overall comparison – Highlighted values shows the best model with algorithms

Model	Accuracy IBM	AUC IBM	Accuracy Occupancy	AUC Occupancy
<u>SVM</u>	<u>89.31%</u>	<u>0.7486354</u>	<u>97%</u>	<u>0.9855</u>
<u>Decision Tree</u>	<u>84.54%</u>	<u>0.5893736</u>	<u>97.86%</u>	<u>0.9825</u>
<u>Boosting</u>	<u>87.04%</u>	<u>0.6497767</u>	<u>96.55%</u>	<u>0.9874</u>
<u>KNN</u>	<u>84.49%</u>	<u>0.5443645</u>	<u>99.41%</u>	<u>0.9937284</u>
<u>ANN</u>	<u>89.37%</u>	<u>0.8964721</u>	<u>91.23%</u>	<u>0.94877</u>

References:

<https://gist.github.com/primaryobjects/d02b93f1e539a9dd2c85>,
http://topepo.github.io/caret/train-models-by-tag.html#Neural_Network
<https://beckmw.wordpress.com/tag/nnet/>